

```
import chess
import random

# Define the Chess Environment
board = chess.Board()

# Define Rewards
WIN_REWARD = 1
LOSS_REWARD = -1
MOVE_REWARD = 0.1

# Create a Q-Table
q_table = {}

# Training Loop
num_episodes = 100
learning_rate = 0.1
discount_factor = 0.9
epsilon = 0.1

for episode in range(num_episodes):
    state = board.fen() # Current state of the game
    total_reward = 0

    while not board.is_game_over():
        # Exploration vs. Exploitation
        if random.uniform(0, 1) < epsilon:
            # Random action
            action = random.choice(list(board.legal_moves))
        else:
            # Greedy action based on Q-values
```

```

if state not in q_table:
    q_table[state] = {}
legal_moves = list(board.legal_moves)
q_values = [q_table[state].get(move.uci(), 0) for move in legal_moves]
max_q = max(q_values, default=0)
best_moves = [move for move, q in zip(legal_moves, q_values) if q == max_q]
action = random.choice(best_moves)

# Take action and observe the next state and reward
move = chess.Move.from_uci(action.uci())
board.push(move)
next_state = board.fen()
reward = MOVE_REWARD

if board.is_checkmate():
    # Win the game
    reward = WIN_REWARD
elif board.is_stalemate() or board.is_insufficient_material() or board.is_seventyfive_moves():
    # Draw the game
    reward = 0

total_reward += reward

# Q-learning update
if state not in q_table:
    q_table[state] = {}
if next_state not in q_table:
    q_table[next_state] = {}

old_q = q_table[state].get(action.uci(), 0)
max_next_q = max(q_table[next_state].values(), default=0)

```

```

new_q = (1 - learning_rate) * old_q + learning_rate * (reward + discount_factor * max_next_q)
q_table[state][action.uci()] = new_q

state = next_state

# Print episode statistics
print(f"Episode: {episode+1}, Total Reward: {total_reward}")

# Reset the chessboard for the next episode
board.reset()

# Testing
num_test_games = 10
wins = 0

for _ in range(num_test_games):
    board.reset() # Reset the board before each test game
    while not board.is_game_over():
        state = board.fen()
        if state not in q_table:
            # If state is not in Q-table, choose a random move
            move = random.choice(list(board.legal_moves))
        else:
            # Choose the best move based on Q-values
            legal_moves = list(board.legal_moves)
            q_values = [q_table[state].get(move.uci(), 0) for move in legal_moves]
            max_q = max(q_values, default=0)
            best_moves = [move for move, q in zip(legal_moves, q_values) if q == max_q]
            move = random.choice(best_moves)
        move = chess.Move.from_uci(move.uci())
        board.push(move)

```

```
if board.is_checkmate():  
    wins += 1  
  
win_rate = wins / num_test_games  
print(f"Win rate against random play: {win_rate}")
```

Episode: 1, Total Reward: 51.300000000000046
Episode: 2, Total Reward: 40.700000000000031
Episode: 3, Total Reward: 20.600000000000023
Episode: 4, Total Reward: 32.600000000000019
Episode: 5, Total Reward: 29.400000000000148
Episode: 6, Total Reward: 40.600000000000031
Episode: 7, Total Reward: 33.500000000000206
Episode: 8, Total Reward: 32.900000000000002
Episode: 9, Total Reward: 34.800000000000225
Episode: 10, Total Reward: 44.700000000000365
Episode: 11, Total Reward: 22.300000000000047
Episode: 12, Total Reward: 49.700000000000436
Episode: 13, Total Reward: 36.600000000000025
Episode: 14, Total Reward: 33.400000000000205
Episode: 15, Total Reward: 20.100000000000016
Episode: 16, Total Reward: 39.800000000000296
Episode: 17, Total Reward: 47.600000000000406
Episode: 18, Total Reward: 38.800000000000028
Episode: 19, Total Reward: 27.800000000000125
Episode: 20, Total Reward: 43.200000000000344
Episode: 21, Total Reward: 46.900000000000004
Episode: 22, Total Reward: 50.800000000000045
Episode: 23, Total Reward: 30.100000000000158
Episode: 24, Total Reward: 40.200000000000003
Episode: 25, Total Reward: 32.300000000000019
Episode: 26, Total Reward: 24.900000000000084
Episode: 27, Total Reward: 15.099999999999966
Episode: 28, Total Reward: 34.200000000000216
Episode: 29, Total Reward: 51.200000000000046
Episode: 30, Total Reward: 49.800000000000044
Episode: 31, Total Reward: 57.700000000000055
Episode: 32, Total Reward: 26.500000000000107
Episode: 33, Total Reward: 33.700000000000021
Episode: 34, Total Reward: 40.000000000000003
Episode: 35, Total Reward: 30.000000000000156
Episode: 36, Total Reward: 38.600000000000028
Episode: 37, Total Reward: 46.400000000000039
Episode: 38, Total Reward: 32.800000000000196
Episode: 39, Total Reward: 37.000000000000256
Episode: 40, Total Reward: 33.900000000000021
Episode: 41, Total Reward: 43.000000000000034
Episode: 42, Total Reward: 42.000000000000033
Episode: 43, Total Reward: 37.900000000000027
Episode: 44, Total Reward: 67.200000000000046
Episode: 45, Total Reward: 58.100000000000556
Episode: 46, Total Reward: 37.700000000000266
Episode: 47, Total Reward: 56.500000000000053
Episode: 48, Total Reward: 58.400000000000056
Episode: 49, Total Reward: 29.700000000000152
Episode: 50, Total Reward: 53.800000000000495
Episode: 51, Total Reward: 29.700000000000152
Episode: 52, Total Reward: 42.300000000000033
Episode: 53, Total Reward: 31.800000000000182
Episode: 54, Total Reward: 32.800000000000196
Episode: 55, Total Reward: 20.900000000000013
Episode: 56, Total Reward: 31.800000000000182
Episode: 57, Total Reward: 29.800000000000153
Episode: 58, Total Reward: 11.899999999999977

Episode: 59, Total Reward: 44.300000000000036
Episode: 60, Total Reward: 29.700000000000152
Episode: 61, Total Reward: 42.500000000000334
Episode: 62, Total Reward: 71.60000000000021
Episode: 63, Total Reward: 53.00000000000048
Episode: 64, Total Reward: 26.500000000000107
Episode: 65, Total Reward: 33.80000000000021
Episode: 66, Total Reward: 42.40000000000033
Episode: 67, Total Reward: 26.200000000000102
Episode: 68, Total Reward: 24.400000000000077
Episode: 69, Total Reward: 35.20000000000023
Episode: 70, Total Reward: 49.40000000000043
Episode: 71, Total Reward: 41.40000000000032
Episode: 72, Total Reward: 29.900000000000155
Episode: 73, Total Reward: 41.00000000000031
Episode: 74, Total Reward: 36.900000000000254
Episode: 75, Total Reward: 40.80000000000031
Episode: 76, Total Reward: 35.00000000000023
Episode: 77, Total Reward: 40.30000000000003
Episode: 78, Total Reward: 44.30000000000036
Episode: 79, Total Reward: 27.800000000000125
Episode: 80, Total Reward: 38.70000000000028
Episode: 81, Total Reward: 33.60000000000021
Episode: 82, Total Reward: 18.09999999999973
Episode: 83, Total Reward: 28.200000000000117
Episode: 84, Total Reward: 26.900000000000112
Episode: 85, Total Reward: 29.10000000000013
Episode: 86, Total Reward: 37.90000000000027
Episode: 87, Total Reward: 34.900000000000226
Episode: 88, Total Reward: 32.60000000000019
Episode: 89, Total Reward: 30.100000000000158
Episode: 90, Total Reward: 14.99999999999966
Episode: 91, Total Reward: 40.20000000000003
Episode: 92, Total Reward: 46.000000000000384
Episode: 93, Total Reward: 36.50000000000025
Episode: 94, Total Reward: 42.40000000000033
Episode: 95, Total Reward: 14.09999999999997
Episode: 96, Total Reward: 19.9
Episode: 97, Total Reward: 17.29999999999996
Episode: 98, Total Reward: 33.400000000000205
Episode: 99, Total Reward: 36.80000000000025
Episode: 100, Total Reward: 38.50000000000028
Win rate against random play: 0.9

Aim:

The aim is to train a chess-playing agent using reinforcement learning to win games by making effective moves and learning optimal strategies through interaction with the environment.

Algorithm:

1. Initialize the Q-table to store Q-values for state-action pairs.
2. Define the hyperparameters, such as learning rate and exploration rate.
3. Training loop:
 - a. Initialize the chess game board.
 - b. Set the initial state of the game board.
 - c. Repeat until a maximum number of episodes or convergence condition is reached:
 - i. Choose an action based on the current state and exploration-exploitation trade-off.
 - ii. Execute the chosen action, observe the new state and reward.
 - iii. Update the Q-value of the current state-action pair using the Q-learning update equation.
 - iv. Update the current state to the new state.
 - d. Save the learned Q-table.
4. Testing phase:
 - a. Initialize the chess game board.
 - b. Set the initial state of the game board.
 - c. Repeat a fixed number of test games:
 - i. Choose an action based on the current state and learned Q-values.
 - ii. Execute the chosen action, observe the new state and reward.
 - iii. Update the current state to the new state.
 - iv. Repeat until the game is over.
 - d. Calculate and display the win rate against a randomly playing opponent.
5. Evaluate the agent's performance by analyzing the win rate and other metrics.

Explanation

Certainly! Here's a simplified explanation of the steps involved in implementing reinforcement learning to train a chess-playing agent:

1. Define the Chess Environment: Create a representation of the chessboard and the rules of the game.

2. Define Actions: Determine the possible moves that the agent can make based on the current state of the game.
3. Define Rewards: Design a system to provide feedback to the agent based on its actions, encouraging winning and discouraging losing.
4. Create a Q-Table: Set up a table to store the expected rewards for each state-action pair.
5. Training Loop: Repeat the following steps for multiple episodes (complete games):
 - a. Select an action based on the current state and the Q-table.
 - b. Observe the next state and the reward received after taking the action.
 - c. Update the Q-value for the previous state-action pair based on the observed reward and the maximum Q-value of the next state.
6. Exploration vs. Exploitation: Balance between trying out new actions and exploiting the learned knowledge in the Q-table.
7. Training Termination: Decide when to stop training, either after a specific number of episodes or when the agent reaches a desired level of performance.
8. Testing: Evaluate the performance of the trained agent by playing against human players or other chess engines.

It's important to note that for complex games like chess, traditional Q-learning with a tabular Q-table may not be feasible. Advanced methods like deep Q-learning or policy gradient methods are often used to handle the game's complexity by approximating the Q-values or policy using neural networks.

Program

```
import chess
import random

# Define the Chess Environment
board = chess.Board()

# Define Rewards
WIN_REWARD = 1
LOSS_REWARD = -1
MOVE_REWARD = 0.1

# Create a Q-Table
q_table = {}

# Training Loop
num_episodes = 10
learning_rate = 0.1
discount_factor = 0.9
epsilon = 0.1

for episode in range(num_episodes):
    state = board.fen() # Current state of the game
    total_reward = 0

    while not board.is_game_over():
        # Exploration vs. Exploitation
        if random.uniform(0, 1) < epsilon:
            # Random action
            action = random.choice(list(board.legal_moves))
        else:
```

```

# Greedy action based on Q-values
if state not in q_table:
    q_table[state] = {}
legal_moves = list(board.legal_moves)
q_values = [q_table[state].get(move.uci(), 0) for move in legal_moves]
max_q = max(q_values, default=0)
best_moves = [move for move, q in zip(legal_moves, q_values) if q == max_q]
action = random.choice(best_moves)

# Take action and observe the next state and reward
move = chess.Move.from_uci(action.uci())
board.push(move)
next_state = board.fen()
reward = MOVE_REWARD

if board.is_checkmate():
    # Win the game
    reward = WIN_REWARD
elif board.is_stalemate() or board.is_insufficient_material() or board.is_seventyfive_moves():
    # Draw the game
    reward = 0

total_reward += reward

# Q-learning update
if state not in q_table:
    q_table[state] = {}
if next_state not in q_table:
    q_table[next_state] = {}

old_q = q_table[state].get(action.uci(), 0)

```

```

max_next_q = max(q_table[next_state].values(), default=0)
new_q = (1 - learning_rate) * old_q + learning_rate * (reward + discount_factor * max_next_q)
q_table[state][action.uci()] = new_q

state = next_state

# Print episode statistics
print(f"Episode: {episode+1}, Total Reward: {total_reward}")

# Reset the chessboard for the next episode
board.reset()

# Testing
num_test_games = 10
wins = 0

for _ in range(num_test_games):
    board.reset() # Reset the board before each test game
    while not board.is_game_over():
        state = board.fen()
        if state not in q_table:
            # If state is not in Q-table, choose a random move
            move = random.choice(list(board.legal_moves))
        else:
            # Choose the best move based on Q-values
            legal_moves = list(board.legal_moves)
            q_values = [q_table[state].get(move.uci(), 0) for move in legal_moves]
            max_q = max(q_values, default=0)
            best_moves = [move for move, q in zip(legal_moves, q_values) if q == max_q]
            move = random.choice(best_moves)
        move = chess.Move.from_uci(move.uci())

```

```
board.push(move)

if board.is_checkmate():
    wins += 1

win_rate = wins / num_test_games
print(f"Win rate against random play: {win_rate}")
```

Output:

```
Episode: 1, Total Reward: 25.500000000000092
Episode: 2, Total Reward: 39.200000000000029
Episode: 3, Total Reward: 37.300000000000026
Episode: 4, Total Reward: 11.099999999999998
Episode: 5, Total Reward: 38.7000000000000266
Episode: 6, Total Reward: 34.000000000000021
Episode: 7, Total Reward: 41.9000000000000325
Episode: 8, Total Reward: 17.399999999999963
Episode: 9, Total Reward: 37.400000000000026
Episode: 10, Total Reward: 55.800000000000052
Win rate against random play: 0.6
```

Result:

